

OneLP

The AI-native operating system for Limited Partners

Turning fragmented private-market documents into structured intelligence, predictive insight, and automated workflows

Kishan Dhulashia · Max Vanderlinden · Sofia Mollon · Agustin Gentil

Venture Capstone · IE School of Science & Technology · 2026



Kishan Dhulashia

Co-founder


[role / background — TEAM]



Max Vanderlinden

Co-founder

[role / background — TEAM]



Sofia Mollon

Co-founder

[role / background — TEAM]



Agustin Gentil

Co-founder

[role / background — TEAM]

IE School of Science & Technology — MSc in Business Analytics & Big Data. Combined background across data science, software engineering, and private markets.

PROBLEM

LPs run multi-million-euro portfolios on PDFs and email

- Private-market portfolios are managed across fragmented systems — GP portals, PDFs, emails, spreadsheets — with no standard data schema.

Consequences for every LP and family office:

- 30–40 hours/month of manual reporting per family office
- Missed capital calls → late-payment penalties & compliance risk
- No real-time view of exposure, concentration, or liquidity
- 5+ disconnected tools to assemble one portfolio picture

The real problem is not aggregation — it is operationalizing unstructured, multi-source financial data into real-time, actionable intelligence.

30–40h

manual reporting / month
per family office

5+

fragmented tools
to replace

16+

GP document formats
no standard schema

Market opportunity

TEAM TO COMPLETE

TAM [€ — TEAM]

Global private-markets / LP software spend

SAM [€ — TEAM]

EU family offices, MFOs & LPs (target segments)

SOM [€ — TEAM]

Reachable in 3 yrs via pilot network & GTM

Sizing methodology (TEAM to complete):

- TAM = # LPs/family offices globally × annual software & data spend
- SAM = EU SFO + MFO + LP segments OneLP can serve today
- SOM = realistic 3-yr capture given pilot pipeline & GTM capacity

Evidence of demand (have):

- 7 active pilots representing €5B+ aggregate AUM
- Family-office wealth growing; private-markets allocations rising
- Sources to cite: Preqin / Campden Wealth / Deloitte FO reports

From fragmented documents to decisions

1 · INGEST

- AI document pipeline
- 16+ doc types
- OCR + LLM hybrid extraction
- Field-level evidence & audit trail

2 · UNDERSTAND

- Real-time NAV / IRR / TVPI / DPI
- Exposure & concentration
- Monte Carlo cash-flow forecasting
- Stress tests & liquidity reserves

3 · ACT

- Capital-call alerts & tracking
- AI daily briefing
- Scheduled reports
- AI email drafting to GPs

The wedge: an **LP interpretation & decision layer** above generic monitoring — not 'we store your data', but 'we turn reports, notices and LPAs into verified answers, alerts and decisions'.

Competitive landscape & differentiation

Player	Category	Strong at	Gap OneLP exploits
eFront / iLEVEL	Enterprise platforms	Breadth, institutional scale	Heavy, not LP-opinionated
Chronograph / Cobalt	Monitoring + reporting	Dashboards, doc-aware AI	Monitoring, not decisions
Canoe	Doc extraction layer	Collection & extraction	Plumbing, not intelligence
OneLP	LP intelligence layer	Interpret · forecast · automate	EU-ready, explainable, fast

Our four defensible wedges:

- (1) LP interpretation layer — fees, waterfall, valuation, 'what changed' in plain language
- (2) Decision workflows — pacing, liquidity, commitment recommendations (not just dashboards)
- (3) Europe-ready compliance — AIFMD / SFDR / DORA built in, not bolted on
- (4) Explainable, source-linked AI + faster time-to-value for lean LP teams

Live product demonstration

► LIVE DEMO

Demo flow (≈4 min):

1. Upload a real GP capital-call PDF → watch the pipeline classify & extract
2. Review queue: click-to-highlight evidence on each extracted field
3. Auto-apply to fund records → NAV / metrics / capital-call notice update live
4. Dashboard: exposure, liquidity, concentration update in real time
5. Forecasting tab: Monte Carlo bands, stress tests, daily AI briefing

Live at the deployed platform · demo accounts seeded (Huergo FO, F3 Finance MFO, Atlas MFO). Fallback: recorded screen-capture if connectivity fails.

Document extraction workflow

1 · DETECT

Two-stage classification: filename/context heuristic (early-exit ≥ 0.70) → Gemini LLM across 16+ doc types

2 · EXTRACT

Document AI (tables/forms, bounding boxes) + LLM (narrative) run in PARALLEL

3 · MERGE

Precedence-based conflict resolution (parser > 0.85 wins, else LLM); conflicts logged

4 · VALIDATE

Schema + per-type confidence thresholds (0.70–0.95); low-confidence fields flagged

5 · APPLY

High-confidence → auto-apply to fund records; else → human review queue

Every field carries raw text, confidence, page number and bounding box → full auditability and click-to-highlight verification.

System architecture

BUILT & VERIFIED

APPLICATION Dashboard · review queue · daily briefing · workflow management

WORKFLOW AUTOMATION Async task queue · scheduled exports · alert engine · AI email drafting


ANALYTICS Monte Carlo · stress testing · liquidity reserve · J-curve profiles · explainability

PROCESSING / INGESTION Classification · hybrid extraction · merge · validation · auto-apply

STORAGE PostgreSQL (structured) · Cloud Storage (documents) · cache (forecasts)

Cloud-native on Google Cloud · multi-tenant isolation (clientId) · field-level auditability · RBAC (LP read-only) · async processing · Sentry observability.

ML/DS #1 — extraction accuracy (hybrid beats either engine)


**0.968**Hybrid macro F1 target ≥ 0.90 ✓**0.968**Critical-field F1 target ≥ 0.95 ✓**0.889**

Doc-type accuracy 8 types


ML/DS #2 — forecasting & backtest (MAE 5.0% < 15%)

BUILT & VERIFIED


Portfolio Net Cash-Flow Forecast — Monte Carlo Confidence Bands
(1,000 paths, independent paths*, base case)



Backtest — Forecast vs Actual (most-recent origin shown)
Rolling MAE \$0.30M = 5.0% of gross flow (20 pts); 80% band coverage 65%



ML/DS #3 — calibration, approach rationale & lessons



Why these choices:

- Hybrid extraction — parser excels at tables, LLM at narrative; merge captures both
- Monte Carlo over closed-form — captures tail risk & cross-asset correlation
- Per-type confidence thresholds — capital calls need 0.85, reports tolerate 0.70

Challenges & lessons learned:

- Pipeline is conservatively calibrated (ECE 0.056) — under-states reliability → thresholds can be safely relaxed
- Correlation matrix goes singular on multi-fund books → engine falls back to independent simulation (documented; fix = regularize matrix)
- End-to-end latency P50 24s / P95 29s — parallel extraction is the key design choice (Metric 3 met)

Traction & market validation

7

active pilots

SFOs · MFOs · LPs

€5B+

aggregate AUM

across pilots

Top 3

of 140+ startups

IE Venture Day

- Live, deployed product with active users iterating in feedback loops
- Live pilot deals: F3 Finance (Belgian MFO) · BEMO Group (SFO) — signed pilot terms

Institutional support:

- Google Cloud for Startups — accepted & funded (infrastructure & architecture costs)
- LvlUp Labs Venture Accelerator — selected
- Best Startup — IE School of Science & Technology
- Demo accounts span real client archetypes (Huergo FO, Atlas MFO, F3 Finance)

Business model — six-tier SaaS pricing

Tier	Annual (EUR)	Users	Funds
UHNWI	€5,000	1	15
SFO Core	€15,000	5	25
SFO White Glove	€30,000	10	75
MFO Platform	€50,000	15+	unlimited
MFO White Glove	€75–100K	unlimited	unlimited
Enterprise	€100–150K	unlimited	unlimited

- Annual SaaS subscription by client type
- Add-ons: seats (€1.5–2K), funds (€200), family workspaces (€4–6K)
- Pilot program: 6 mo free + infra cost, converts to paid
- Selling now: UHNWI, SFO Core, MFO Platform
- White Glove / Enterprise open 2028+

Unit economics & 5-year financials

	2026	2027	2028	2029	2030
Net revenue	€8K	€110K	€556K	€2.09M	€4.94M
Clients (EOY)	7	22	50	91	141
Gross margin	80%	91%	94%	95%	96%
EBIT	-€73K	-€239K	-€330K	-€89K	+€1.23M
LTV : CAC	14.2x	37.7x	23.7x	20.0x	35.2x
CAC payback	4 mo	4 mo	5 mo	7 mo	5 mo

€4.94M

2030 net revenue
from €8K in 2026

>14x

LTV:CAC every year
capital-efficient

25%

2030 EBIT margin
profitable at scale

Go-to-market strategy

TEAM TO COMPLETE

Target customers (have):

- Single & multi-family offices, LPs, advisory teams — EU-first
- Beachhead: lean MFO/SFO teams drowning in GP paperwork

Motion so far (have):

- Pilot partner program → warm intros to FO network
- Founder-led sales; bi-weekly pilot touchpoints

Scale-up GTM (TEAM to complete):

- Channel strategy: outbound, partnerships, FO networks, events
- Sales cycle length & conversion by tier (fill from pilot data)
- Geographic expansion sequence beyond EU
- CAC by channel & payback targets
- Referral / network-effect loops

Product roadmap

NOW

- Ingestion pipeline (16+ types)
- Dashboard & analytics
- Monte Carlo forecasting
- Daily briefing & alerts
- 7 pilots live

NEXT (6–12 mo)

- Fee / waterfall / carry verification
- ILPA v2.0 performance logic
- Pacing & commitment engine
- EU compliance (AIFMD/SFDR/DORA)
- Convert pilots → paid

LATER (12–24 mo)

- Portfolio-company KPI pages
- Benchmarking & PME
- White-Glove & Enterprise tiers
- Geographic expansion
- [TEAM: prioritize]

Funding requirements & use of funds

Raising: €300K pre-seed

23 months runway to prove product-market fit

Then: €750K seed (mid-2027)

→ self-sustaining profitability by 2030 with €1.4M in the bank

Total dilutive capital to breakeven: ~€1M

Use of pre-seed funds:

- Product development & engineering
- First 3 hires (eng + customer success)
- Pilot client onboarding & conversion
- Cloud infra (partly covered by Google Cloud credits)

Monthly burn: €3K (2026 H1) → €20K (2027) → cash-positive 2030. **TEAM: confirm the current ask amount, valuation, and instrument (SAFE / equity) before pitching.**

OneLP

The data-driven operating system for private markets

Live product · 7 pilots · €5B+ AUM · evaluated to F1 0.968, forecast MAE 5.0%, latency P50 24s
Built on real AI, validated with real clients, measured against real targets.

Thank you · Questions?

Kishan · Max · Sofia · Agustin | info@onelp.capital